

Recommender System

Exam Project of *“Advanced Information Management 2”*

FEDERICO MOTTA

Università degli studi di Modena e Reggio Emilia

April 23, 2020

Outline

■ Introduction

- *Problem definition*
- *State of the Art*

■ GRank

- *Tripartite Preference Graph*
- *Algorithm*
- *Time complexity*
- *Memory voracity*

■ Naïve Bayes

- *Vector Space Model*
- *Algorithm*
- *Standard Natural Language Preprocessing*

■ Results

- *Performances*
- *NDCG Plot*
- *Accuracy & Execution Time Plots*

■ Software

- *Architecture & Implementation choices*
- *Output Example*

■ Conclusions

Problem definition

Input:

k = *number of items to suggest*

$\mathbf{U} = \{u_1, \dots, u_M\}$

users

$\mathbf{I} = \{i_1, \dots, i_N\}$

items

$\mathbf{O} = \{\langle u, i, j \rangle : \text{"user } u \text{ prefers item } i \text{ over item } j \text{"}\}$

observations

Note: all type of user preferences (e.g. rating, like, browsing history) can be converted to $\langle u, i, j \rangle$ observations

Output:

given a user \tilde{u} return k never seen items which he probably appreciates

Papers



Shams, B., & Haratizadeh, S. (2016). Graph-based collaborative ranking. *arXiv preprint arXiv:1604.03147*. Available at <https://arxiv.org/abs/1604.03147>.



Lops, P., de Gemmis, M., & Semeraro, G. (2011). Content-based recommender systems: State of the art and trends. In F. Ricci, L. Rokach, B. Shapira, & P. B. Kantor (Eds.), *Recommender systems handbook* (pp. 73–105). doi:10.1007/978-0-387-85820-3_3




Basu, C., Hirsh, H., & Cohen, W. (1998). Recommendation as classification: Using social and content-based information in recommendation. In *Proceedings of the national conference on artificial intelligence (aaai '98)*. Available at <http://cs.cmu.edu/~wcohen/postscript/aaai-98-collab.ps>, AAAI Press.



Baeza-Yates, R., Ribeiro-Neto, B. et al. (1999). Text operations. In *Modern information retrieval* (Chap. 7, Vol. 463, pp. 163–190, with Ziviani, Nivio). Available at <https://homepages.dcc.ufmg.br/~nivio/papers/aw99.ps>. Addison–Wesley.


Collaborative-filtering (GRank) (Shams & Haratizadeh, 2016)

Tripartite Preference Graph



Collaborative-filtering (GRank) (Shams & Haratizadeh, 2016)

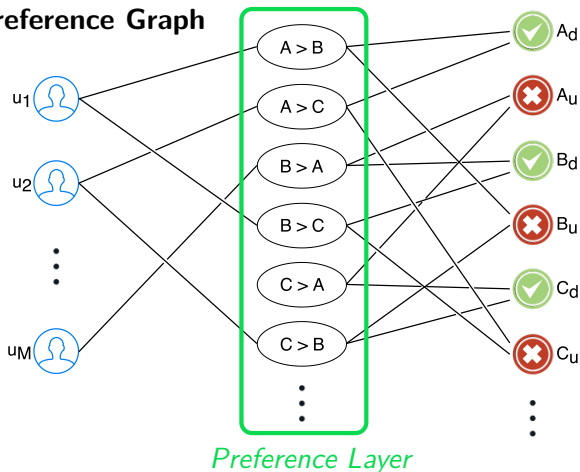
Tripartite Preference Graph



Representative Layer

Collaborative-filtering (GRank) (Shams & Haratizadeh, 2016)

Tripartite Preference Graph



Notation:

$$\mathbf{P} = \{p_1, \dots, p_{N \cdot (N-1)}\}$$

$$p = \langle i, j \rangle \in \mathbf{P} = "i > j"$$

preferences

"item i is preferred to item j "

Ranking Algorithm

$$\mathbf{GR} : U \times I \rightarrow \mathbb{R}_{\geq 0}$$

$$(1) \quad GR(i) = \frac{\overline{PPR}_{[i_d]}}{\overline{PPR}_{[i_d]} + \overline{PPR}_{[i_u]}}$$

“goodness” of unseen item i

Author's note: overlined symbols are vectors

Ranking Algorithm

$$\mathbf{GR} : U \times I \rightarrow \mathbb{R}_{\geq 0}$$

$$(1) \quad GR(i) = \frac{\overline{PPR}_{[i_d]}}{\overline{PPR}_{[i_d]} + \overline{PPR}_{[i_u]}}$$

“goodness” of unseen item i

$$(2) \quad \overline{PPR}_t = \alpha \cdot T \cdot \overline{PPR}_{t-1} + (1 - \alpha) \cdot \overline{PV}$$

Personalized Page Rank

Author's note: overlined symbols are vectors

Ranking Algorithm

$$\mathbf{GR} : U \times I \rightarrow \mathbb{R}_{\geq 0}$$

$$(1) \quad GR(i) = \frac{\overline{PPR}_{[i_d]}}{\overline{PPR}_{[i_d]} + \overline{PPR}_{[i_u]}}$$

“goodness” of unseen item i

$$(2) \quad \overline{PPR}_t = \alpha \cdot T \cdot \overline{PPR}_{t-1} + (1 - \alpha) \cdot \overline{PV}$$

Personalized Page Rank

$$(3) \quad T_{[ij]} = \begin{cases} \frac{1}{\text{degree}(i)} & \text{if } (i, j) \in E \\ 0 & \text{otherwise} \end{cases}$$

Transition Matrix

$$(4) \quad PV_{[j]} = \begin{cases} 1 & \text{if } j = \tilde{u} \\ 0 & \text{otherwise} \end{cases}$$

Personalized Vector

Author's note: overlined symbols are vectors

Grank Complexity

$$\blacksquare \quad |V| = \underbrace{M}_{\text{users}} + \overbrace{N \cdot (N - 1)}^{\text{preferences}} + \underbrace{2 \cdot N}_{\text{representatives}}$$

$$\blacksquare \quad |E| = \underbrace{S}_{\text{user-preference}} + \overbrace{2 \cdot N \cdot (N - 1)}^{\text{preference-representative}}$$

$$S \simeq c \cdot N^2 \quad c \approx 2.48$$

Grank Complexity

$$\blacksquare \quad |V| = \underbrace{M}_{\text{users}} + \overbrace{N \cdot (N - 1)}^{\text{preferences}} + \underbrace{2 \cdot N}_{\text{representatives}}$$

$$\blacksquare \quad |E| = \underbrace{S}_{\text{user-preference}} + \overbrace{2 \cdot N \cdot (N - 1)}^{\text{preference-representative}}$$

$$S \simeq c \cdot N^2 \quad c \approx 2.48$$

$$\blacksquare \quad \text{graph construction} \simeq O(N^2 + M)$$

space complexity

$$\blacksquare \quad \text{recommendation} \simeq O(t \cdot E) = O(N^2)$$

time complexity

$$t \leq 20$$

Grank Complexity

$$\blacksquare \quad |V| = \underbrace{M}_{\text{users}} + \overbrace{N \cdot (N - 1)}^{\text{preferences}} + \underbrace{2 \cdot N}_{\text{representatives}}$$

$$\blacksquare \quad |E| = \underbrace{S}_{\text{user-preference}} + \overbrace{2 \cdot N \cdot (N - 1)}^{\text{preference-representative}} \quad S \simeq c \cdot N^2 \quad c \approx 2.48$$

$$\blacksquare \quad \text{graph construction} \simeq O(N^2 + M) \quad \text{space complexity}$$

$$\blacksquare \quad \text{recommendation} \simeq O(t \cdot E) = O(N^2) \quad \text{time complexity} \quad t \leq 20$$

Datasets built with $\mathbb{T} = 30, 40, 50, 60$ reviews^a for each user ($M = 100, 300, 1000$).

^a using a *stratified random sampling* as suggested in “Basu, Hirsh, and Cohen, 1998”

Transition Matrix

Alert: building the transition matrix with $\mathbb{T} = 30$ and $M \simeq 2\,000$ takes **26 hours** and **11.4 GB** of RAM (using a CSR matrix)

Transition Matrix

Alert: building the transition matrix with $\mathbb{T} = 30$ and $M \simeq 2000$ takes **26 hours** and **11.4 GB** of RAM (using a CSR matrix) **that is too much !**

Transition Matrix

Alert: building the transition matrix with $\mathbb{T} = 30$ and $M \simeq 2000$ takes **26 hours** and **11.4 GB** of RAM (using a CSR matrix) **that is too much !**

$$T = \begin{matrix} & \begin{matrix} \text{users} & \text{preferences} & \text{representatives} \end{matrix} \\ \begin{matrix} \text{users} \\ \text{preferences} \\ \text{representatives} \end{matrix} & \begin{pmatrix} - & - & 1 & - & - & - \\ 2 & - & - & - & - & 3 \\ - & - & + & - & + & - \\ - & - & - & 5 & 6 & - \end{pmatrix} \end{matrix}$$

■ density of T is $\approx 2.4 \cdot 10^{-8}$

Transition Matrix

Alert: building the transition matrix with $\mathbb{T} = 30$ and $M \simeq 2000$ takes **26 hours** and **11.4 GB** of RAM (using a CSR matrix) **that is too much !**

$$T = \begin{matrix} & \begin{matrix} \text{users} & \text{preferences} & \text{representatives} \end{matrix} \\ \begin{matrix} \text{users} \\ \text{preferences} \\ \text{representatives} \end{matrix} & \begin{pmatrix} & & & & \\ & 1 & & & \\ - & 2 & - & & 3 \\ - & - & + & - & \\ - & - & & 5 & 6 & - \end{pmatrix} \end{matrix}$$

The matrix T is a 3x3 block matrix. The top row is labeled 'users', the middle row 'preferences', and the bottom row 'representatives'. The columns are also labeled 'users', 'preferences', and 'representatives'. The matrix is partitioned into blocks: block 1 (top-left), block 2 (top-middle), block 3 (top-right), block 4 (middle-left), block 5 (middle-middle), block 6 (middle-right), block 7 (bottom-left), block 8 (bottom-middle), and block 9 (bottom-right). Block 4 is highlighted with a red box.

- density of T is $\approx 2.4 \cdot 10^{-8}$
- this is the **largest** submatrix
- all elements in block 4 are $1/2$
- all elements in block 6 are $1/(n-1)$

Transition Matrix

Alert: building the transition matrix with $\mathbb{T} = 30$ and $M \simeq 2000$ takes **26 hours** and **11.4 GB** of RAM (using a CSR matrix) **that is too much !**

$$T = \begin{matrix} & \begin{matrix} \text{users} & \text{preferences} & \text{representatives} \end{matrix} \\ \begin{matrix} \text{users} \\ \text{preferences} \\ \text{representatives} \end{matrix} & \begin{pmatrix} & & & & \\ & 1 & & & \\ - & 2 & - & & 3 \\ - & - & + & - & \\ - & - & - & 4 & \\ & 5 & 6 & & \end{pmatrix} \end{matrix}$$

The matrix T is a 3x3 block matrix. The top-left block (1,1) contains the value 1. The top-right block (1,3) contains the value 3. The middle-left block (2,1) contains the value 2. The middle-right block (2,3) contains the value 4. The bottom-left block (3,1) contains the value 5. The bottom-right block (3,3) contains the value 6. The block containing 4 and 6 is highlighted with a red box.

- density of T is $\approx 2.4 \cdot 10^{-8}$
- this is the **largest** submatrix
- all elements in block 4 are $1/2$
- all elements in block 6 are $1/(n-1)$

Thus, T has been split in:

$$T_1 = \text{blocks } 1, 2, 3, 5$$

$$T_2 = \text{blocks } 4 \text{ and } 6 \quad (\text{generated row by row with the } \textit{yield} \text{ statement})$$

Since $T = T_1 + T_2$ nothing changes a part performances ✓

Keyword-based (Naïve Bayes)^(Lops, de Gemmis, & Semeraro, 2011)

Vector Space Model:

$$\mathbf{D} = \{ d_1, \dots, d_N \}$$

corpus (of documents)

$$\mathbf{T} = \{ t_1, \dots, t_Z \}$$

dictionary (of keywords/terms)

$$\mathbf{d}_j = \{ w_{1j}, \dots, w_{Zj} \} \in D$$

weights of terms $t_1 \dots t_Z$ in document j

Keyword-based (Naïve Bayes)^(Lops et al., 2011)

Vector Space Model:

$$\mathbf{D} = \{ d_1, \dots, d_N \}$$

corpus (of documents)

$$\mathbf{T} = \{ t_1, \dots, t_Z \}$$

dictionary (of keywords/terms)

$$\mathbf{d}_j = \{ w_{1j}, \dots, w_{Zj} \} \in D$$

weights of terms $t_1 \dots t_Z$ in document j

Weighting the terms:

TF-IDF handles well $\overbrace{\text{single/multiple occurrences}}^{\text{TF assumption}}, \overbrace{\text{rare/frequent terms}}^{\text{IDF assumption}}$

$$\text{TF-IDF}(t_k, d_j) = \frac{\overbrace{f_{k,j}}^{\text{Term Frequency}}}{\max_z f_{z,j}} \cdot \log \underbrace{\frac{N}{n_k}}_{\text{Inverse Document Frequency}}$$

$$n_k = |\{ d \in D \text{ with at least one } t_k \}|$$

$$f_{k,j} = \text{frequency of } t_k \text{ in } d_j$$

Keyword-based (Naïve Bayes)^(Lops et al., 2011)

Vector Space Model:

$$\begin{aligned}
 \mathbf{D} &= \{ d_1, \dots, d_N \} && \text{corpus (of documents)} \\
 \mathbf{T} &= \{ t_1, \dots, t_Z \} && \text{dictionary (of keywords/terms)} \\
 \mathbf{d}_j &= \{ w_{1j}, \dots, w_{Zj} \} \in D && \text{weights of terms } t_1 \cdots t_Z \text{ in document } j
 \end{aligned}$$

Weighting the terms:

TF-IDF handles well $\overbrace{\text{single/multiple occurrences}}^{\text{TF assumption}}$, $\overbrace{\text{rare/frequent terms}}^{\text{IDF assumption}}$, $\overbrace{\text{long/short docs.}}^{\text{normalization assumption}}$.

$$\text{TF-IDF}(t_k, d_j) = \underbrace{\frac{f_{k,j}}{\max_z f_{z,j}}}_{\text{Term Frequency}} \cdot \underbrace{\log \frac{N}{n_k}}_{\text{Inverse Document Frequency}}$$

$n_k = |\{d \in D \text{ with at least one } t_k\}|$
 $f_{k,j} = \text{frequency of } t_k \text{ in } d_j$

$$\mathbf{w}_{k,j} = \text{TF-IDF}(t_k, d_j) / \sqrt{\sum_{s=1}^{|\mathbf{T}|} \text{TF-IDF}(t_s, d_j)^2} \quad (\text{cosine normalization})$$

Ranking Algorithm

$$\mathbf{C}_{\text{MAP}} = \underbrace{\arg \max_{c_j} P(c_j/d_i)}_{\substack{\text{maximum a posteriori probability} \\ \text{of document } d_i \text{ belonging to class } c_j \\ \text{(most likely class)}}} = \arg \max_{c_j} \frac{P(c_j) \cdot P(d_i/c_j)}{\cancel{P(d_i)}}$$

Ranking Algorithm

$$\mathbf{C}_{\text{MAP}} = \underbrace{\arg \max_{c_j} P(c_j/d_i)}_{\substack{\text{maximum a posteriori probability} \\ \text{of document } d_i \text{ belonging to class } c_j \\ \text{(most likely class)}}} = \arg \max_{c_j} \frac{P(c_j) \cdot P(d_i/c_j)}{\cancel{P(d_i)}}$$

Assumptions:

- bag of words
- conditional independence

Ranking Algorithm

$$\begin{aligned}
 \mathbf{C}_{\text{MAP}} &= \underbrace{\arg \max_{c_j} P(c_j/d_i)}_{\substack{\text{maximum a posteriori probability} \\ \text{of document } d_i \text{ belonging to class } c_j \\ \text{(most likely class)}}} = \arg \max_{c_j} \frac{P(c_j) \cdot P(d_i/c_j)}{\cancel{P(d_i)}}
 \end{aligned}$$

Assumptions:

- bag of words
- conditional independence

Multinomial Naïve Bayes Model:

$$P(c_j/d_i) = P(c_j) \cdot \prod_{w \in d_i} P(w/c_j)$$

Descriptions Preprocessing^b

- unescape HTML entities (`"&"` → `"&"`)
- substitute punctuation with spaces (`"requirements:minimum"` → `"requirements minimum"`)
- remove multiple spaces (`"Washington DC"` → `"Washington" "DC"`)
- split camelCase words (`"GoldenEye"` → `"Golden Eye"`)
- lowercase (`"Xbox"` → `"xbox"`)

^b Baeza-Yates, Ribeiro-Neto, et al., 1999 ~ Chapter 7, "Text Operations"

Descriptions Preprocessing^b

- unescape HTML entities (`"&"` → `"&"`)
- substitute punctuation with spaces (`"requirements:minimum"` → `"requirements minimum"`)
- remove multiple spaces (`"Washington DC"` → `"Washington" "DC"`)
- split camelCase words (`"GoldenEye"` → `"Golden Eye"`)
- lowercase (`"Xbox"` → `"xbox"`)
- tokenization (`"Frisco and LA"` → `"Frisco" "and" "LA"`)

^b Baeza-Yates, Ribeiro-Neto, et al., 1999 ~ Chapter 7, "Text Operations"

Descriptions Preprocessing^b

- unescape HTML entities (`"&"` → `"&"`)
- substitute punctuation with spaces (`"requirements:minimum"` → `"requirements minimum"`)
- remove multiple spaces (`"Washington DC"` → `"Washington" "DC"`)
- split camelCase words (`"GoldenEye"` → `"Golden Eye"`)
- lowercase (`"Xbox"` → `"xbox"`)
- tokenization (`"Frisco and LA"` → `"Frisco" "and" "LA"`)
- stop words removal (`"drive" "a" "train"` → `"drive" "train"`)

^b Baeza-Yates, Ribeiro-Neto, et al., 1999 ~ Chapter 7, "Text Operations"

Descriptions Preprocessing^b

- unescape HTML entities (`"&"` → `"&"`)
- substitute punctuation with spaces (`"requirements:minimum"` → `"requirements minimum"`)
- remove multiple spaces (`"Washington DC"` → `"Washington" "DC"`)
- split camelCase words (`"GoldenEye"` → `"Golden Eye"`)
- lowercase (`"Xbox"` → `"xbox"`)
- tokenization (`"Frisco and LA"` → `"Frisco" "and" "LA"`)
- stop words removal (`"drive" "a" "train"` → `"drive" "train"`)
- stemming (with Porter) (`"playable"` → `"play"`)

^b Baeza-Yates, Ribeiro-Neto, et al., 1999 ~ Chapter 7, "Text Operations"

Descriptions Preprocessing^b

- unescape HTML entities (`"&"` → `"&"`)
- substitute punctuation with spaces (`"requirements:minimum"` → `"requirements minimum"`)
- remove multiple spaces (`"Washington DC"` → `"Washington" "DC"`)
- split camelCase words (`"GoldenEye"` → `"Golden Eye"`)
- lowercase (`"Xbox"` → `"xbox"`)
- tokenization (`"Frisco and LA"` → `"Frisco" "and" "LA"`)
- stop words removal (`"drive" "a" "train"` → `"drive" "train"`)
- stemming (with Porter) (`"playable"` → `"play"`)
- lemmatization (with WordNet) (optional)

^b Baeza-Yates, Ribeiro-Neto, et al., 1999 ~ Chapter 7, "Text Operations"

Performances

Normalized Discounted Cumulative Gain is a standard strategy used to assess recommender systems by comparing the quality of their top-k suggestions.

Performances

Normalized Discounted Cumulative Gain is a standard strategy used to assess recommender systems by comparing the quality of their top- k suggestions.

$$\text{NDCG@k} = \frac{\sum_{i=1}^k \frac{\overbrace{2^{r_i^u} - 1}^{\text{rating of user } u \text{ to } i\text{-th item}}}{\log_2(i+1)}}{\underbrace{\sum_{i=1}^k \frac{\overbrace{2^5 - 1}^{\text{max rating is 5 } \star}}{\log_2(i+1)}}_{\text{normalization coefficient} = \alpha_u}}$$

Performances

Normalized Discounted Cumulative Gain is a standard strategy used to assess recommender systems by comparing the quality of their top- k suggestions.

$$\text{NDCG@}k = \frac{\sum_{i=1}^k \frac{\overbrace{2^{r_i^u} - 1}^{\text{rating of user } u \text{ to } i\text{-th item}}}{\log_2(i+1)}}{\underbrace{\sum_{i=1}^k \frac{2^5 - 1}{\log_2(i+1)}}_{\text{max rating is 5 } \star}} \quad \text{normalization coefficient} = \alpha_u$$


Other useful metrics are:

- accuracy = $\frac{TP+TN}{TP+FP+FN+TN}$
- precision = $\frac{TP}{TP+FP}$
- recall = $\frac{TP}{TP+FN}$
- F_1 score = $\frac{2 \cdot \text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$


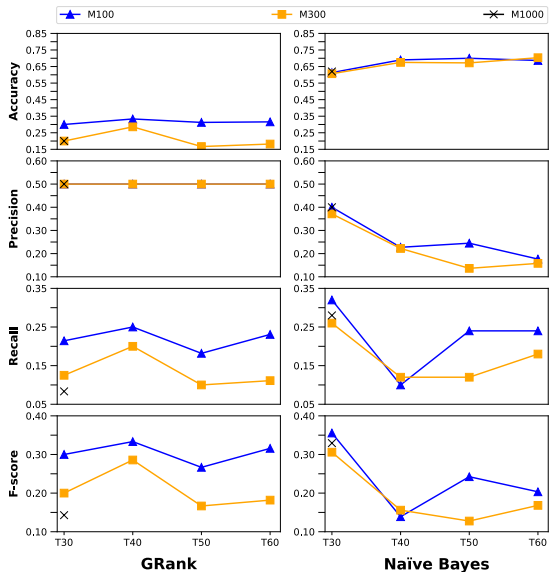
	real positive	real negative
predicted positive	TP	FP
predicted negative	FN	TN

confusion matrix

NDCG performance comparison



Other performance comparison



Development choices (source code available at <https://github.com/TheJena/RecommenderSystem>)

- Shrink MONGODB size

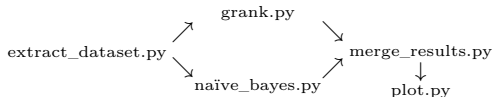
*(instructions in: **docs/how_to_create_test_db_dump.pdf**)*

Development choices (source code available at <https://github.com/TheJena/RecommenderSystem>)

- Shrink MONGODB size

*(instructions in: **docs/how_to_create_test_db_dump.pdf**)*

- Modular Architecture



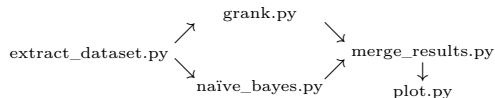
- Extract data from MONGODB once, then load a JSON file

Development choices (source code available at <https://github.com/TheJena/RecommenderSystem>)

- Shrink MONGODB size

(instructions in: **docs/how_to_create_test_db_dump.pdf**)

- Modular Architecture



- Extract data from MONGODB once, then load a JSON file
- ~~NETWORKX~~ → SCIPY + generator-iterators^c + multiprocessing^d

^c <https://docs.python.org/3/glossary.html#term-generator-iterator>

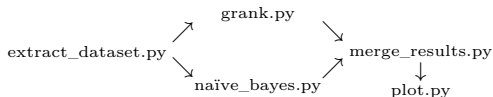
^d <https://docs.python.org/3/library/multiprocessing.html>

Development choices (source code available at <https://github.com/TheJena/RecommenderSystem>)

- Shrink MONGODB size

(instructions in: [docs/how_to_create_test_db_dump.pdf](#))

- Modular Architecture



- Extract data from MONGODB once, then load a JSON file
- ~~NETWORKX~~ → SCIPY + generator-iterators^c + multiprocessing^d
- NATURAL LANGUAGE TOOL KIT + SCI KIT LEARN + WordNet Expansion
- (in content-based) rating normalization, scaling (e.g. $\times 1000$) and rounding, fit, sort, use top-quartile^e as a threshold to update confusion matrix

^c <https://docs.python.org/3/glossary.html#term-generator-iterator>

^d <https://docs.python.org/3/library/multiprocessing.html>

^e "Basu et al., 1998"

Output Example

How to run, from a BASH shell:

```
./extract_dataset.py --category "Books" --category "Video Games" --output dataset/M100_T30.json
                        --users      100      --reviews  30                                --threads 8

./grank.py      dataset/M100_T30.json --stop-after 50 -k 10 --output results/grank.yaml      --parallel 8
./naive_bayes.py dataset/M100_T30.json --stop-after 50 -k 10 --output results/naive_bayes.yaml

./merge_results.py --content-based results/naive_bayes.yaml --stop-after 1
                  --graph-based  results/grank.yaml         --top-k      10
```

Sample output:

===== Recommendations for user "A11L3YX5WIDKJ" =====

```
naive bayes weight coefficient: 0.097      <~ NDCG@k
grank weight coefficient: 0.490            <~ NDCG@k
```

position	item	weighted rating
1	B005FYJA52	0.430
2	0060788380	0.364
3	0007147295	0.335
4	B0038MTE7C	0.333
5	0006550436	0.331
...

Conclusions

GRank

✓ NDCG similar to paper results

Multinomial Naïve Bayes

✗ worst NDCG performance

Conclusions

GRank

- ✓ NDCG similar to paper results
- ✓ higher F_1 score

Multinomial Naïve Bayes

- ✗ worst NDCG performance
- ✗ lower precision

Conclusions

GRank

- ✓ NDCG similar to paper results
- ✓ higher F_1 score
- ✗ slower (days)

Multinomial Naïve Bayes

- ✗ worst NDCG performance
- ✗ lower precision
- ✓ much faster (hours)

Conclusions

GRank

- ✓ NDCG similar to paper results
- ✓ higher F_1 score
- ✗ slower (days)
- ✗ scales $\propto N^2$ (n° items)

Multinomial Naïve Bayes

- ✗ worst NDCG performance
- ✗ lower precision
- ✓ much faster (hours)
- ✓ scales $\propto N$ (n° items)

Conclusions

GRank

- ✓ NDCG similar to paper results
- ✓ higher F_1 score
- ✗ slower (days)
- ✗ scales $\propto N^2$ (n° items)
- ✓ scales $\propto M$ (n° users)

Multinomial Naïve Bayes

- ✗ worst NDCG performance
- ✗ lower precision
- ✓ much faster (hours)
- ✓ scales $\propto N$ (n° items)

Conclusions

GRank

- ✓ NDCG similar to paper results
- ✓ higher F_1 score
- ✗ slower (days)
- ✗ scales $\propto N^2$ (n° items)
- ✓ scales $\propto M$ (n° users)
- ✓ scales $\propto S$ (n° user's preferences)

Multinomial Naïve Bayes

- ✗ worst NDCG performance
- ✗ lower precision
- ✓ much faster (hours)
- ✓ scales $\propto N$ (n° items)